

# Multicore Processors: Architecture & Programming

## Lab 1

In this first lab you will implement a histogram in OpenMP.

### What is the problem definition?

Your program reads a text file that consists of a random combination of the four letters: *a*, *b*, *c*, and *d*. The file can contain any number of these characters and not in a specific order. Your program must determine the which of the four letters occurred the most and print the result on the screen.

For example: *b occurred the most 400 times of a total of 1000 characters.*

The output means that the letter b occurred the most frequently among the four letters and it occurred 400 times out of the 1000 characters that were in the file.

### What is the input?

The input to your program consists of three things:

- N: number of threads, can be 0 (purely sequential), 1 (OpenMP version with one thread), or 4 (OpenMP with four threads).
- num: number of characters in the file
- Filename: This is the name of the file that contains the characters. The format of the file is simple. It consists of a character followed by another character, etc. For example:  
acabdaddd

We are providing you with a file random-char.c to generate the characters file. Feel free to use that file, compile and run it to generate the files, or build your own file generator.

Assume your program is called *maxnum*, the command line is expected to be:

```
./maxnum N num filename
```

You must generate the filename first, either by using random-char.c or by hand if you want, before executing the above command.

### What is output?

Your program must print on the screen:

*x occurred the most y times of a total of z characters.*

Where x is the character that occurred the most (a, b, c, or d) and y is the number of times x has occurred in the file and z is the total number of characters in the file.

## How will you solve this?

Here is an algorithm for a sequential version:

```
Assume we have an array of four integers, one for each of the four characters
while(not done reading the characters)
{
    Read a character
    Determine the entry of the array to be incremented
    Increment the entry
}
Scan the four entries of the array and determine the highest one
Print the output on the screen
```

For the OpenMP version, the one with four threads, you need to do the following (Note: Some details are left on purpose for you to figure out):

- Load the whole file in an array of characters (i.e. string). Assume its name is: list[]
- Each thread will be responsible for  $N/4$  characters. You need to take care of the corner case when  $N$  is not divisible by 4.
- Assume  $N/4 = \text{step}$
- Thread 0 will take care of characters from 0 to step-1. Thread 1 will take care of characters from step to  $2 \cdot \text{step} - 1$ , and so one.
- Let each thread have its own array of four int where it puts the count of the four characters.
- At the end, combine the four arrays into one.
- Finally, pick the largest one and print the output.

## Refresh Your Knowledge about C

OpenMP is built on top of C in our case. To write the program, you need to refresh your memory about the following C items. This list is not exhaustive but contains the most important items and the ones that people tend to forget if they don't use the language a lot.

- How to read inputs from the command line. This has to do with the arguments of the main() function: argc and argv.
- How to open/close a file.
- How to use fread()
- How to use printf
- Dynamic allocation with malloc() and free()

## The report

Beside your source code, you need to submit a report based that contains the results of several experiments as will be discussed. You will use the **time** command in Linux. You

need to use the same CIMS machine (crunchy1, crunchy3, ...) for all your runs. The **time** command, as we saw in class, produces three numbers: real, usr, sys.

You need to fill out two tables.

Table 1 contains the *real* part of the **time** command.

N →	1,000	10,000	100,000	1,000,000	10,000,000	100,000,000
Seq						
1-thread						
4-threads						

The sequential is a purely sequential code. The 1-thread contains the OpenMP commands but generates 1 thread. The 4-threads is OpenMP with 4 threads. The first row shows the number of characters in different files. So, you need to generate those different files.

Table 2 contains the *cpu+sys* part of the **time** command.

N →	1,000	10,000	100,000	1,000,000	10,000,000	100,000,000
Seq						
1-thread						
4-threads						

Draw two bar-graphs. Both of them have N as the x-axis.

- In graph 1: The y-axis is the speedup (time seq / time 4 threads) where time is the *real* of **time** command. That is, you will get the numbers from the first table.
- In graph 2: The y-axis is the efficiency = speedup (that you calculated in first graph)/4

Finally answer the following questions at the end of the report:

1. Which machine have you used to do the experiment (crunchy1, crunchy3, ...)? Doing them on your laptop will not be accepted.
2. Given the Table 1 and graph 1 what is the relationship between N and the speedup?
3. How can you justify that relationship you showed in the above question?
4. Given Table 1 and Table 2: Was the *real* time mostly bigger, smaller, equal to *cpu+sys* time?
5. What is your justification for the trend in question 4 above?
6. What is the trend you see in graph 2?
7. How do you explain the trend you found in question 6 above?

### Regarding compilation & naming convention:

- Name your source file: netID-code.c (where netID is your netID number).
- Name your report netID-report.pdf (where netID is your netID number).
- Do the compilation and execution on crunchyx (x = 1, 3, 5, or 6) machines.
- Use the latest version of gcc on crunchy by typing: **module load gcc-9.2** (or higher)
- Compile with: **gcc -fopenmp -Wall -std=c99 -o maxnum netID-code.c**

**What to submit?**

Add netID-code.c and netID-report.pdf to a single zip file netID.zip

**How to submit?** Through the assignment sections of Brightspace.

**Have Fun!**