

Multicore Processors: Architecture & Programming

Lab# 2

In this second lab you will write parallel code, on CIMS, crunchy, machines, to generate prime numbers from 2 to N and test scalability and performance.

General notes:

- Program will be written in OpenMP.
- Load gcc-12.2 from crunchy using: `module load gcc-12.2`
- The name of the source code file is: `genprime.c`
- You compile it as: `gcc -O3 -std=c99 -Wall -fopenmp -o genprime genprimes.c -lm`
- Write your program in such a way that to execute it I will type: `./genprime N t`
Where :
N is a positive number bigger than 2
t is the number of threads and is a positive integer that does not exceed 100.
- The output of your program is a text file `N.txt` (N is the number entered as argument to your program).
- Assume we will not do any tricks with the input (i.e. We will not deliberately test your program with wrong values of N or t).

The format of the output file `N.txt`

- one prime per line
- each line has the format: a b
 - a: the rank of the number (1 means the first prime)
 - b: the number itself
- Assume the first line of the file to be: 1 2
- The second line will then be: 2 3
- and the third: 3 5
- and so on.

The algorithm for generating prime numbers:

There are many algorithms for generating prime numbers and for primality testing. Some are more efficient than others. For this lab, we will implement the following algorithm, given N:

1. Generate all numbers from 2 to N.
2. First number is 2, so remove all numbers that are multiple of 2 (i.e. 4, 6, 8, ... N). Do not remove the 2 itself.
3. Following number is 3, so remove all multiple of 3 that have not been removed from the previous step. That will be: 9, 15, ... till you reach N.
4. The next number that has not been crossed so far is 5. So, remove all multiple of 5 that have not been crossed before, till you reach N.
5. Continue like this till $\text{floor}((N+1)/2)$.
6. The remaining numbers are the prime numbers.

Example:

Suppose $N = 20$

floor of $(20+1)/2 = 10$ ← where we stop.

Initially we have:

2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20

Let's cross all multiple of 2 (but leave 2):

2, 3, 4, 5, ~~6~~, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, 15, ~~16~~, 17, ~~18~~, 19, ~~20~~

Next number is 3, so we cross all multiple of 3 that have not been crossed:

2, 3, 4, 5, ~~6~~, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, ~~15~~, ~~16~~, 17, ~~18~~, 19, ~~20~~

Next number that has not been crossed is 5, so we will cross multiple of 5 (i.e. 10, 15, and 20).
As you see below, they are all already crossed.

2, 3, 4, 5, ~~6~~, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, ~~15~~, ~~16~~, 17, ~~18~~, 19, ~~20~~

Next number that has not been crossed is 7, so we will cross multiple of 7 (i.e. 14). As you see below, they are all already crossed.

2, 3, 4, 5, ~~6~~, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, ~~15~~, ~~16~~, 17, ~~18~~, 19, ~~20~~

The next number that has not been crossed is 11. This is bigger than 10, so we stop here.

The numbers that have not been crossed are the prime numbers:

2, 3, 5, 7, 11, 13, 17, 19

The file that your program generates is 20.txt and looks like:

1, 2
2, 3
3, 5
4, 7, 2
5, 11
6, 13
7, 17
8, 19

How to measure the execution time?

Your code will be doing three major tasks:

- Read the input from the command line.
- Generate the prime numbers (as indicated by the algorithm above).
- Write the output file and exit.

We care only about the timing of the middle part. Therefore, you do the following:

```
double tstart = 0.0, tend=0.0, ttaken;
```

Read the input from the command line

```
tstart = omp_get_wtime();
```

Generate the prime numbers (as indicated by the algorithm above) ← This will be the parallel part

```
ttaken = omp_get_wtime() - t_start;
```

```
printf("Time take for the main part: %f\n", ttaken);
```

Write the output file and exit.

The report

Speedup will be calculated using the time measurement indicated above.

- After you implement the OpenMP version of the above algorithm, generate the following graphs:
 - graph 1 (N = 1000), y-axis is the speedup relative to the running time with 1 thread; and x-axis is the number of threads: 2, 5, 10, 25, 50, and 100.
 - graph 2 (N = 1000,000), y-axis is the speedup relative to the running time with 1 thread; and x-axis is the number of threads: 2, 5, 10, 25, 50, and 100.
- For each graph, explain the behavior that you see. Do NOT explain what the curve looks like but why the curve looks the way it is.

Important

- Do not try to change the algorithm. There are other algorithms but, for this lab, we need to implement the one explained above.
- Think carefully about what your data structure will be.
- Implement a sequential version first. This will make your life easier.
- Do not try to produce an optimized parallel code at once. First: sequential code, second: parallel code that is correct, and third: optimized parallel code.
- As a way to help you check the correctness of your code, we are providing you a list of the first 2 million prime numbers in two files. These files are not in the required format explained above. They are just provided for convenience.
- We will not test your code with N bigger than 100 million.

What do you have to submit:

A single zip file. The file name is your netID.zip

Inside that zip file you need to have:

- genprime.c
- pdf file containing the graph and explanation.

The zip file is submitted through Brightspace as usual.

Enjoy!