

LAB 2 discussion

Simulation based Process Scheduling

- Explore multiple Schedulers and what OS tracks
- You will learn how to write a Discrete Event Simulator
- Specification of Process and Random Behavior

Process-1	0	200	40	90
	40	100	10	40
	50	20	10	10
Process-4	60	200	5	20

process
arrival/start

Total
CPU-time

CPU
Burst

IO
Burst

- Cpu-burst = random (range [1.. Proc.cpu_burst])
- Ioburst = random (range [1.. Proc.io_burst])

What you will generate ?

```
$ cat input4
```

```
0 200 40 90
40 100 10 40
50 20 10 10
60 200 5 20
```

```
$ ./sched -sF input4 rfile
```

```
FCFS
```

```
0000: 0 200 40 90 2 | 575 575 360 15
0001: 40 100 10 40 4 | 532 492 300 92
0002: 50 20 10 10 2 | 82 32 10 2
0003: 60 200 5 20 4 | 1174 1114 773 141
SUM: 1174 44.29 81.26 553.25 62.50 0.341
```

↑
Total Sim time

↑
CPU Utilization

↑
I/O Utilization

↑
Avg
Turnaround time

↑
Avg
Wait time

↑
Throughput

static-prio

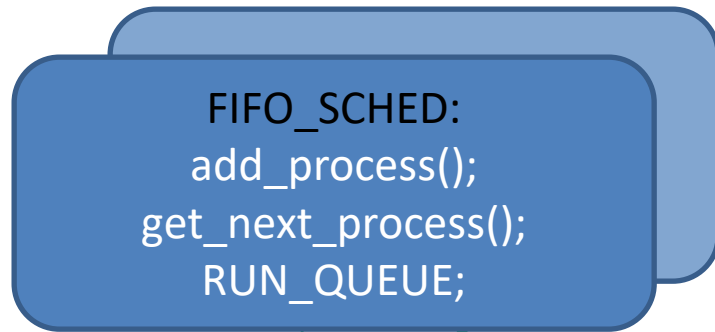
Finish Time

Turnaround time

I/O Time

CPU Wait time

Generic Structure of your Lab-2 infrastructure



Generic Interface
Allows for plugging
Different schedulers

(*add_process());

(*get_next_process());

Simulation and Scheduler Interaction

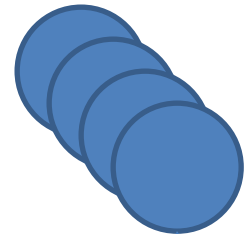
```
While (get_event()) {  
    decode_event() and run action(); // see next page  
}
```

get_event()

put_event()

rm_event()

Processes



Event

Timestamp
Process
oldstate
newstate

DES-Layer



Time ordered

```

void Simulation() {
    EVENT* evt;
    while( (evt = get_event()) ) {
        Process *proc = evt->evtProcess;    // this is the process the event works on
        CURRENT_TIME = evt->evtTimeStamp;
        int transition = evt->transition;    // for accounting
        int timeInPrevState = CURRENT_TIME - proc->state_ts;
        delete evt; evt = nullptr; // remove cur event obj and don't touch anymore

        switch (transition) { // encodes where we come from and where we go
            case TRANS_TO_READY:
                // must come from BLOCKED or CREATED
                // must add to run queue, no event created
                CALL_SCHEDULER = true;
                break;
            case TRANS_TO_PREEMPT: // similar to TRANS_TO_READY
                // must come from RUNNING (preemption)
                // must add to run queue, no event created
                CALL_SCHEDULER = true;
                break;
            case TRANS_TO_RUN:
                // create event for either preemption or blocking
                break;
            case TRANS_TO_BLOCK:
                //create an event for when process becomes READY again
                CALL_SCHEDULER = true;
                break;
        }

        if(CALL_SCHEDULER) {
            if (get_next_event_time() == CURRENT_TIME)
                continue; //process next event from Event queue
            CALL_SCHEDULER = false; // reset global flag
            if (CURRENT_RUNNING_PROCESS == nullptr) {
                CURRENT_RUNNING_PROCESS = THE_SCHEDULER->get_next_process();
                if (CURRENT_RUNNING_PROCESS == nullptr)
                    continue;
                // create event to make this process runnable for same time.
            }
        }
    }
}

```

./sched -v -e input_show rfile

ShowEventQ: 0:0 20:1

0 0 0: CREATED -> READY

AddEvent(0:0:RUNNG): 20:1:READY ==> 0:0:RUNNG 20:1:READY

0 0 0: READY -> RUNNG cb=8 rem=100 prio=1

AddEvent(8:0:BLOCK): 20:1:READY ==> 8:0:BLOCK 20:1:READY

8 0 8: RUNNG -> BLOCK ib=2 rem=92

AddEvent(10:0:READY): 20:1:READY ==> 10:0:READY 20:1:READY

10 0 2: BLOCK -> READY

AddEvent(10:0:RUNNG): 20:1:READY ==> 10:0:RUNNG 20:1:READY

10 0 0: READY -> RUNNG cb=10 rem=92 prio=1

AddEvent(20:0:BLOCK): 20:1:READY ==> 20:1:READY 20:0:BLOCK

20 1 0: CREATED -> READY

20 0 10: RUNNG -> BLOCK ib=7 rem=82

AddEvent(27:0:READY): ==> 27:0:READY

AddEvent(20:1:RUNNG): 27:0:READY ==> 20:1:RUNNG 27:0:READY

20 1 0: READY -> RUNNG cb=7 rem=100 prio=3

AddEvent(27:1:BLOCK): 27:0:READY ==> 27:0:READY 27:1:BLOCK

Input file

0 100 10 10

20 100 20 10

RESULTS OF SIMULATION

FCFS

0000: 0 100 10 10 2 | 234 234 89 45

0001: 20 100 20 10 4 | 226 206 77 29

SUM: 234 85.47 57.26 220.00 37.00 0.855

./sched -v -e input_show rfile

Input file

0 100 10 10
20 100 20 10

ShowEventQ: 0:0 20:1

line triggered by "-v"
Timestamp pid howlong:
FROM -> TO

0 0 0: CREATED -> READY

AddEvent(0:0:RUNNG): 20:1:READY ==> 0:0:RUNNG 20:1:READY

0 0 0: READY -> RUNNG cb=8 rem=100 prio=1

AddEvent(8:0:BLOCK): 20:1:READY ==> 8:0:BLOCK 20:1:READY

One Event

8 0 8: RUNNG -> BLOCK ib=2 rem=92

AddEvent(10:0:READY): 20:1:READY ==> 10:0:READY 20:1:READY

10 0 2: BLOCK -> READY

AddEvent(10:0:RUNNG): 20:1:READY ==> 10:0:RUNNG 20:1:READY

10 0 0: READY -> RUNNG cb=10 rem=92 prio=1

AddEvent(20:0:BLOCK): 20:1:READY ==> 20:1:READY 20:0:BLOCK

line triggered by "-e"
event added: (time,pid,transition)
EventQ-Before → EventQ-After
(should be time ordered)

20 1 0: CREATED -> READY

20 0 10: RUNNG -> BLOCK ib=7 rem=82

AddEvent(27:0:READY): ==> 27:0:READY

AddEvent(20:1:RUNNG): 27:0:READY ==> 20:1:RUNNG 27:0:READY

20 1 0: READY -> RUNNG cb=7 rem=100 prio=3

AddEvent(27:1:BLOCK): 27:0:READY ==> 27:0:READY 27:1:BLOCK

./sched -v -t input_show rfile

Input file

0 100 10 10
20 100 20 10

0 0 0: CREATED -> READY
SCHED (1): 0:0
0 0 0: READY -> RUNNG cb=8 rem=100 prio=1
8 0 8: RUNNG -> BLOCK ib=2 rem=92
SCHED (0):
10 0 2: BLOCK -> READY
SCHED (1): 0:10
10 0 0: READY -> RUNNG cb=10 rem=92 prio=1
20 1 0: CREATED -> READY
20 0 10: RUNNG -> BLOCK ib=7 rem=82
SCHED (1): 1:20
20 1 0: READY -> RUNNG cb=7 rem=100 prio=3
27 0 7: BLOCK -> READY
27 1 7: RUNNG -> BLOCK ib=9 rem=93
SCHED (1): 0:27
27 0 0: READY -> RUNNG cb=7 rem=82 prio=1
34 0 7: RUNNG -> BLOCK ib=1 rem=75
SCHED (0):
35 0 1: BLOCK -> READY
SCHED (1): 0:35
35 0 0: READY -> RUNNG cb=9 rem=75 prio=1
36 1 9: BLOCK -> READY
44 0 9: RUNNG -> BLOCK ib=9 rem=66
SCHED (1): 1:36
44 1 8: READY -> RUNNG cb=11 rem=93 prio=3
53 0 9: BLOCK -> READY
55 1 11: RUNNG -> BLOCK ib=4 rem=82

line triggered by “-t” to show scheduler runqueue

SCHED(len): { “pid:timestamp” }*

where len is length of RQ followed by RQ entries

Current cpu-burst

Current io-burst

Dynamic priority

Remaining cputime

`./sched -p input_show rfile`

To help with the preemption (EPRIO)

There are 2 conditions that need to be met:

Cond1: `unblocked_process.dynprio > current.dynprio` **AND**

Cond2: next event for current proc is > current time

---> PRIO preemption *Cond1*=1 *Cond2*=1 (TS=2) --> YES

---> PRIO preemption *Cond1*=1 *Cond2*=0 (TS=0) --> NO

↓ Time till next event
of current running process

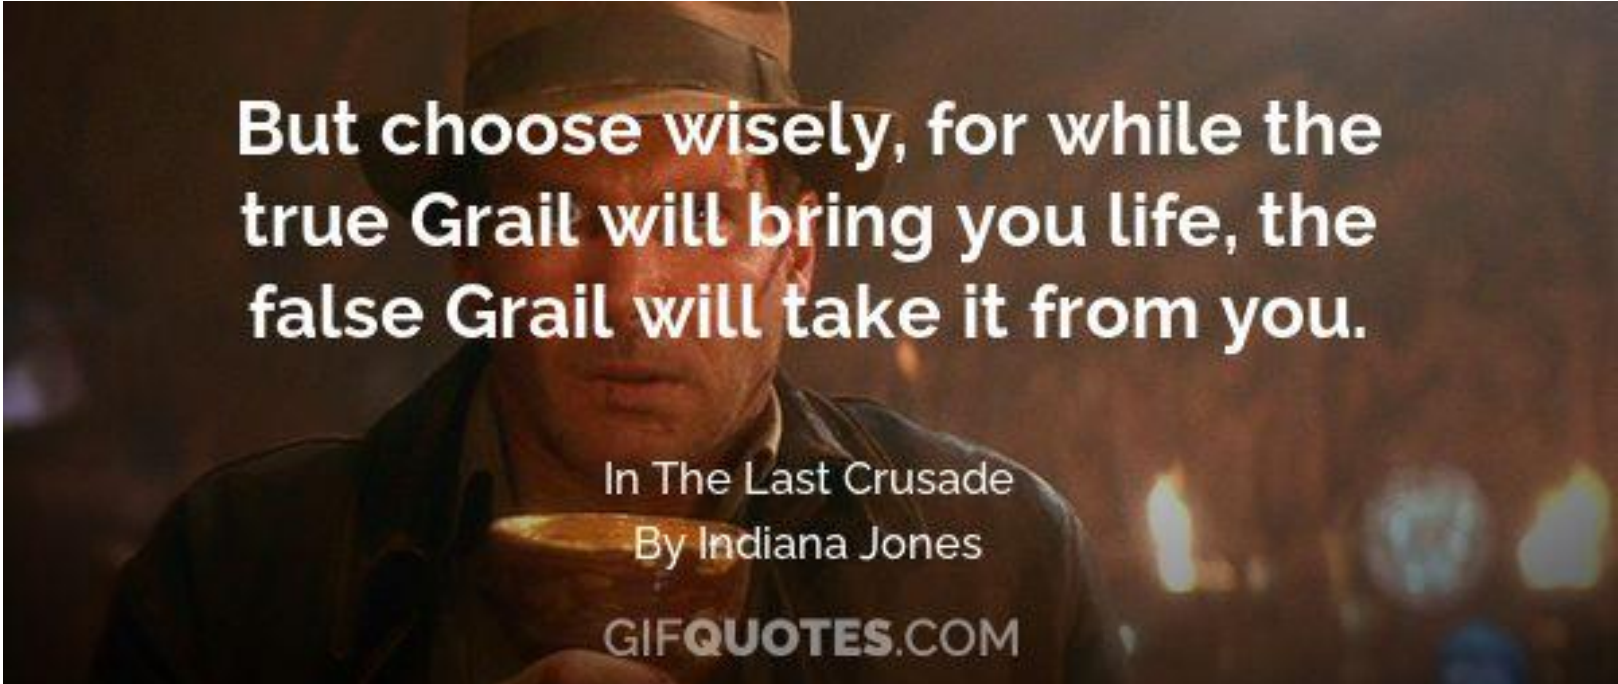
Accounting

- Make sure you do backward accounting (how long was process in a particular state).
- Charge the right bucket (cpu, blocked) when you get out of its state, not when you get in (otherwise "E" scheduler won't work and it forward accounting in real systems is not possible)
- You have all the information how long you were in the state you are transitioning out of.

```
void Simulation() {  
    EVENT* evt;  
    while( (evt = get_event()) ) {  
        Process *proc = evt->evtProcess;    // this is the process the event works on  
        CURRENT_TIME = evt->evtTimeStamp;  
        int transition = evt->transition;    // for accounting  
        int timeInPrevState = CURRENT_TIME - proc->state_ts;  
        delete evt; evt = nullptr; // remove cur event obj and don't touch anymore
```

Data Structures

- Schedulers are often about pooling requests and selecting:
 - Queueing , searching , indexing, mapping, priorities
 - Efficiency matters, not just effectiveness
 - Underlying data structures matter

A still from the movie 'The Last Crusade' showing Indiana Jones (Harrison Ford) wearing his signature fedora and holding a golden chalice (the Grail). The background is dark and atmospheric, with some light flares.

**But choose wisely, for while the
true Grail will bring you life, the
false Grail will take it from you.**

In The Last Crusade
By Indiana Jones

GIFQUOTES.COM

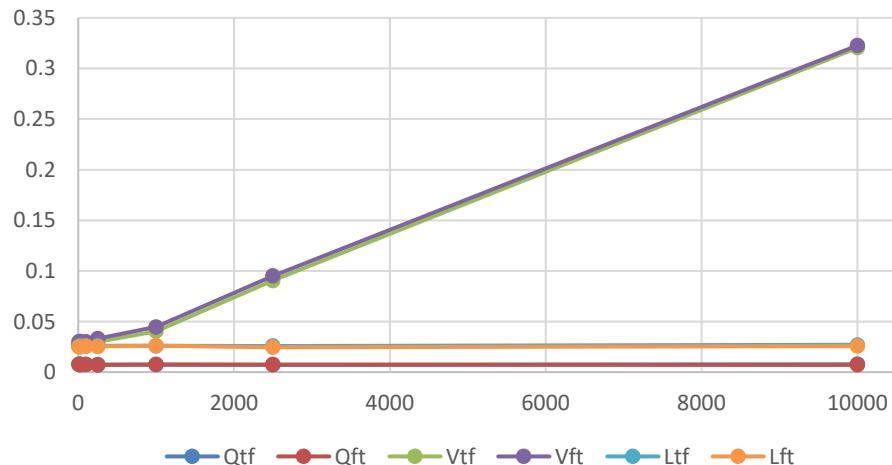
Example: Queue Semantics

- Assume existing queue size N
- Consider the following operations
 - Add front / delete tail (FT)
 - Add Tail / delete front (TF)
 - Iterate through queue with
 - C++ iterator (Iter) or C++ [] operator
- Time these operations
- Consider the following container classes:
 - Deque (Q)
 - Vector (V)
 - List (L)

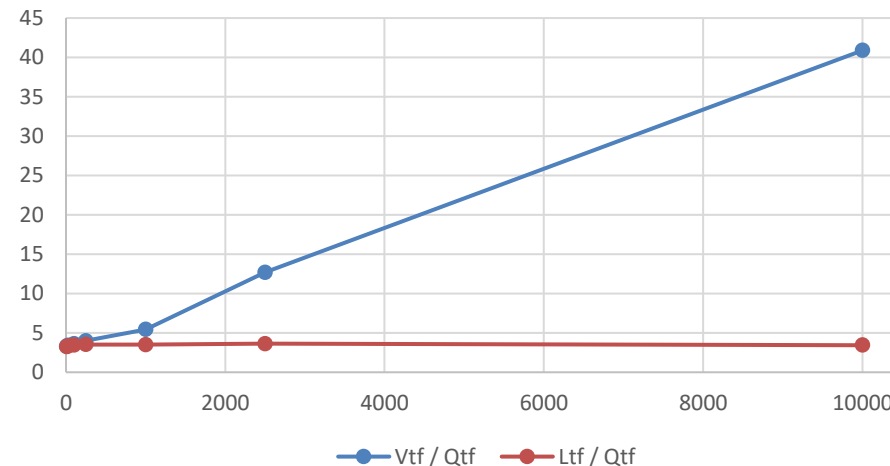
Add / Delete Opposite ends

- Scale the size of the queue (x-axis = N)
- TF and FT semantics
- Why are vectors so expensive ?

Add / Delete at opposite ends



Relative Overhead

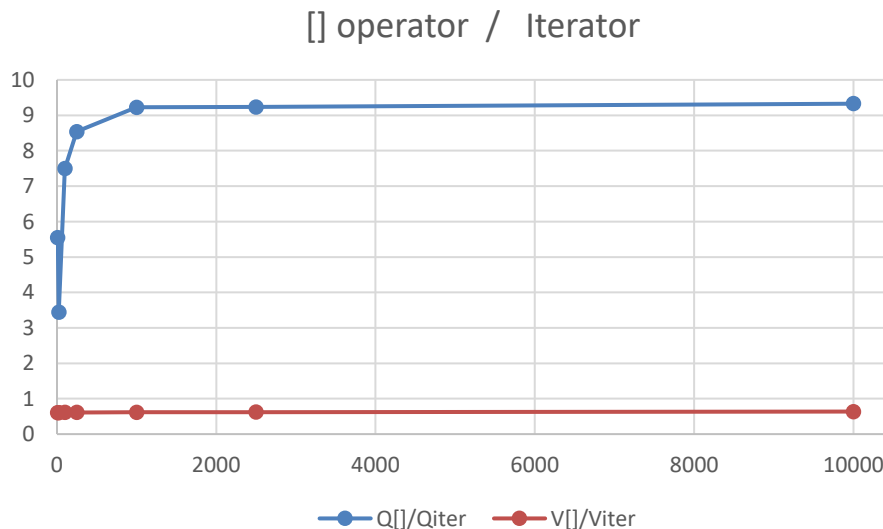


Iterations

- Iterate by 2 approaches:

`iter:` `for (x : objs) { x } ;`

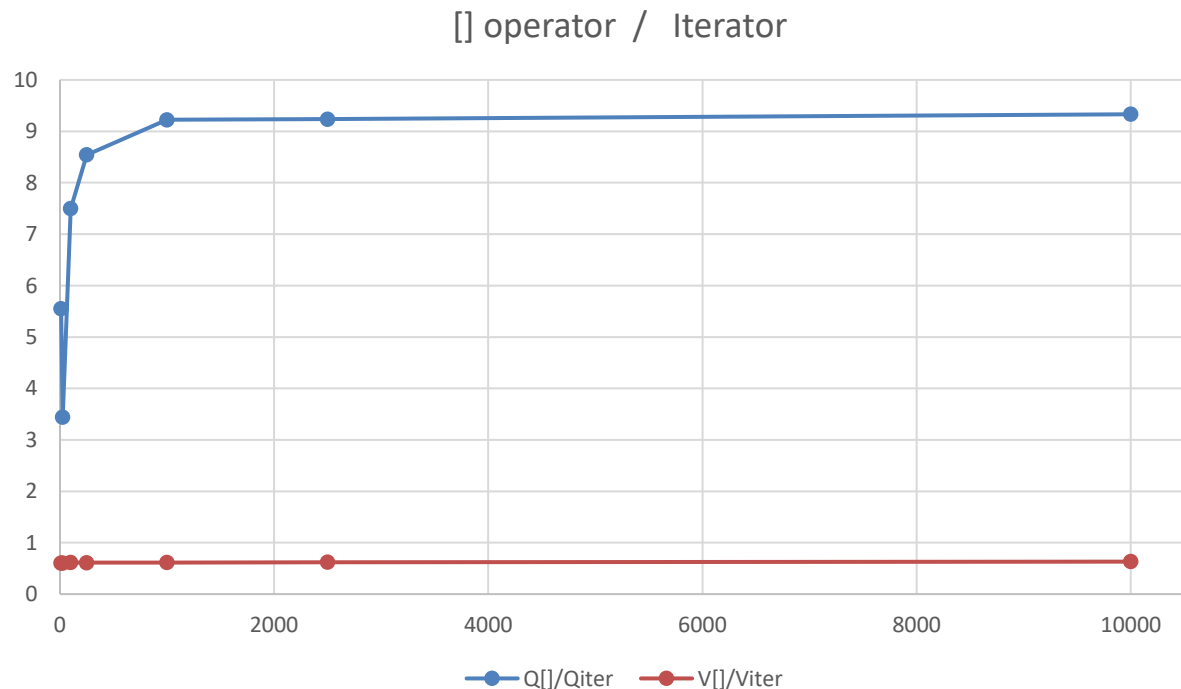
`[]:` `for (i = 0; i<obj.size(); i++) { obj[i]; }`



note: `list[]` does not exist

Iteration vector vs deque


- `[]` operator relatively expensive for deque over iterator, but not for vector



Iteration vector vs deque vs list

- Avoid [] operator on queues
- Think about underlying data structures how something would be implemented and think about cost of operations.

Cost Iterations



Cost Iterations

